

# Gen AI Assignment Questions

## 1. Introduction to Generative AI:

- **Q1:** What is Generative AI? How does it differ from traditional AI models that focus on classification or regression tasks?

**Answer:**

**Generative AI** refers to a class of artificial intelligence models that are designed to generate new content such as text, images, audio, or code. These models learn patterns and structures from training data and can then create similar — yet original — outputs.

Difference between traditional AI and Gen AI

Aspect	Generative AI	Traditional AI (Classification/Regression)
Definition	AI that generates new content like text, images, or audio	AI that makes predictions or classifications based on data
Main Purpose	To create original outputs based on learned patterns	To categorize or predict outcomes from existing data
Output Type	Human-like content (e.g., essays, images, code)	Labels (e.g., spam/ham) or numerical values (e.g., prices)
Example Use Cases	ChatGPT (text), DALL·E (images), Copilot (code)	Spam detection, loan approval, house price prediction
Nature of Output	Creative and variable (many outputs possible)	Deterministic and fixed (one output per input)
Learning Goal	Learn to mimic human-like data creation	Learn to recognize patterns for accurate predictions

- **Q2:** Explain the key differences between **Generative AI** and **Discriminative AI** models. Provide examples of each.

**Answer:**

Aspect	Generative AI	Discriminative AI
<b>Purpose</b>	Learns the <b>joint probability</b> distribution between input and output	Learns the <b>boundary</b> between different classes or outputs
<b>Function</b>	Can <b>generate new data</b> similar to the training data	Focuses on <b>predicting labels or outcomes</b> from input data
<b>Output</b>	Produces content like text, images, audio, etc.	Produces <b>predictions</b> or classifications
<b>Use Case</b>	Content generation, simulation, translation	Email classification, image recognition, fraud detection
<b>Examples</b>	GPT, DALL·E, GANs (Generative Adversarial Networks), VAEs	Logistic Regression, SVM, Decision Trees, Random Forest, BERT
<b>Data Focus</b>	Models <b>how data is generated</b>	Models <b>how to distinguish between outputs</b> based on input
<b>Capabilities</b>	Can do both <b>generation and classification</b>	Can only <b>classify or predict</b>

- **Q3:** Describe the purpose and basic functioning of **Generative Adversarial Networks (GANs)**. What are the roles of the generator and discriminator in a GAN?

**Answer:**

Generative Adversarial Networks (GANs) are a type of **generative AI model** used to create **realistic data** like images, videos, or audio. The goal of a GAN is to **generate new data samples** that are so convincing that they are indistinguishable from real data.

### Basic Functioning of a GAN:

A GAN consists of two neural networks — a Generator and a Discriminator — that are trained together in a game-like setup:

1. **The Generator (G)** tries to create fake data that looks like real data.
2. **The Discriminator (D)** tries to distinguish between real data and fake data produced by the generator.

They compete against each other:

- The generator improves to fool the discriminator.
- The discriminator improves to catch the generator's fakes.

This process continues until the generated data is very close to the real data, and the discriminator can no longer easily tell the difference.

- **Q4: What is a **latent space** in Generative AI? Explain how it is used to generate new data (e.g., images, text).**

#### **Answer:**

Latent space in Generative AI refers to a mathematical representation where data is encoded in a compressed and abstract form. It is a lower-dimensional space that captures the most meaningful features or patterns of the input data, such as images or text, without storing every detail.

In models like GANs (Generative Adversarial Networks) or VAEs (Variational Autoencoders), the original high-dimensional data is mapped into this latent space during training. Each point in the latent space corresponds to a variation of the data, and nearby points often represent similar characteristics.

Once the model is trained, new data can be generated by sampling a random point from this latent space and passing it through the generator (or decoder). The model transforms this abstract representation back into a complete, high-dimensional output, such as a new image or text sequence. For example, generating new faces by sampling different points in the latent space allows the model to produce images that look real but are entirely original.

- **Q5: Explain how **Variational Autoencoders (VAEs)** work. How are VAEs different from GANs, and what types of tasks are VAEs typically used for?**

#### **Answer:**

A Variational Autoencoder (VAE) is a generative model that learns to represent data in a compressed form (latent space) and then reconstructs it. Unlike traditional autoencoders, VAEs learn a probabilistic distribution over the latent space, which allows them to generate new, realistic data by sampling from this distribution.

### **Components of a VAE**

1. **Encoder**

Converts input data into a set of latent variables (mean and standard deviation).

2. **Latent Space**

A compressed space where the model samples points using the encoder's output.

3. **Decoder**

Takes the sampled latent vector and reconstructs the original input.

#### **How VAEs Differ from GANs**

- VAEs use a probabilistic encoder and decoder, while GANs use a generator and a discriminator.
- VAEs aim to reconstruct input data and learn a structured latent space.
- GANs aim to fool a discriminator by generating realistic data.
- VAEs are more stable but may produce blurrier outputs.
- GANs often produce sharper images but can be difficult to train.

#### **Typical Use Cases for VAEs**

- Image generation
- Text synthesis
- Anomaly detection
- Denoising images
- Compressed data representation

## **2. Applications of Generative AI:**

- **Q6:** What are some real-world applications of Generative AI? List at least four areas where generative models are applied and provide a brief description of each.

### **Answer:**

#### **1. Text Generation**

Generative AI models like ChatGPT, Bard, and Claude can generate human-like text. They are used in customer support chatbots, content creation, translation, and summarization. These models understand context and generate meaningful and coherent responses or documents.

#### **2. Image Generation**

Models like DALL-E, Midjourney, and Stable Diffusion can generate realistic or artistic images from text prompts. These are used in graphic design, marketing, game asset creation, and digital art, where original visuals are required quickly.

#### **3. Code Generation**

Tools like GitHub Copilot and CodeWhisperer use generative models to write code based on developer input. These models help with auto-completion, code suggestion, and even building full functions, speeding up software development.

#### 4. Drug Discovery and Healthcare

Generative AI is used in pharmaceuticals to design new molecules with desired properties. It helps in accelerating the discovery of drug candidates, optimizing molecular structures, and simulating protein folding.

- **Q7:** How can Generative AI be applied in the field of healthcare? Provide two examples where it can be useful, such as in drug discovery or medical image generation.

**Answer:**

##### 1. Drug Discovery

Generative AI models can design new molecular structures that have the potential to become effective drugs. These models explore chemical space faster than traditional methods, helping researchers find promising drug candidates with specific properties such as low toxicity or high efficacy. This speeds up the early stages of pharmaceutical research and reduces costs.

##### 2. Medical Image Generation and Enhancement

Generative models like GANs can create or enhance medical images such as X-rays, MRIs, and CT scans. For example, they can generate synthetic medical images to train machine learning models where real data is scarce. They can also improve image resolution or fill in missing parts, which helps radiologists in better diagnosis and treatment planning.

- **Q8:** Explain how **text generation models** (e.g., GPT-3) can be used in content creation. Give an example of how these models can generate blog posts or social media content.

**Answer:**

Text generation models like GPT-3 are designed to produce human-like text based on a given prompt. They are trained on vast amounts of data from books, websites, and articles, allowing them to understand grammar, tone, style, and context. These models are widely used in content creation because they can quickly generate coherent and relevant text across various domains.

**Example: Generating a blog post intro**

*Prompt given to the model:*

"Write an introduction for a blog post on the benefits of remote work."

*Generated Output:*

"Remote work has rapidly become more than just a trend — it's a new way of working that offers flexibility, boosts productivity, and supports a better work-life balance. In this blog post, we'll explore how remote work is transforming modern workplaces and why it could be the future of employment."

#### **Example: Creating a social media post**

*Prompt:*

"Create a tweet promoting a new eco-friendly product line."

*Output:*

"Exciting news! Our brand-new eco-friendly product line is here — made with love and sustainable materials. Shop green, live clean. #EcoFriendly #Sustainability #NewLaunch"

### **3. Popular Generative AI Models:**

- **Q9:** Describe the **GPT** (Generative Pre-trained Transformer) model. How does it generate human-like text, and what are its key applications?

#### **Answer:**

GPT (Generative Pre-trained Transformer) is a type of language model developed by OpenAI that uses the transformer architecture. It is trained on a large amount of text data to learn patterns, grammar, and structure in language. The model is called "pre-trained" because it learns from general text data first and then can be fine-tuned for specific tasks.

#### **How GPT Generates Human-like Text**

GPT generates text using a method called next-word prediction. Given an input prompt, the model predicts the most likely next word based on everything it has seen during training. It repeats this process word by word to form complete sentences and paragraphs. The transformer architecture allows the model to understand context, relationships between words, and long-range dependencies in the input.

#### **Key Applications of GPT**

1. Chatbots and virtual assistants
2. Content generation for blogs, articles, and social media
3. Code generation and software help (e.g., GitHub Copilot)
4. Language translation and summarization

- **Q10:** Explain how **Recurrent Neural Networks (RNNs)** can be used for generating sequences, such as text or music. How do RNNs handle sequential data differently from other models?

**Answer:**

Recurrent Neural Networks (RNNs) are a type of neural network designed to work with sequential data, where the order and context of elements matter. They are commonly used for tasks like text generation, speech recognition, and music composition.

**How RNNs Work with Sequences**

1. They use loops to pass information from one time step to the next.
2. The output at any time depends on both the current input and the previous hidden state.
3. This allows RNNs to model temporal patterns and context over time, unlike standard models that treat each input in isolation.

**Use Cases of RNNs in Sequence Generation**

- **Text generation** (writing sentences, paragraphs)
- **Music composition** (generating melodies or rhythms)
- **Speech synthesis**
- **Time-series forecasting** (e.g., stock prices or weather data)

Unlike traditional feedforward networks that treat each input independently, RNNs have a special architecture that allows them to retain information from previous inputs using hidden states. When processing a sequence, the RNN updates its internal state at each step and uses this state to influence future predictions. This memory of prior inputs makes RNNs suitable for generating sequences where each output depends on the earlier parts.

- **Q11:** What is **BERT** (Bidirectional Encoder Representations from Transformers), and how is it different from models like GPT-3 in the context of text generation?

**Answer:**

BERT is a transformer-based language model developed by Google. It stands for Bidirectional Encoder Representations from Transformers. Unlike traditional models that read text left-to-right or right-to-left, BERT reads text in both directions at once, allowing it to understand the full context of each word in a sentence.

BERT is trained using two main tasks:

1. **Masked Language Modeling** (predicting missing words in a sentence)

2. Next Sentence Prediction (understanding the relationship between two sentences)  
It is mainly used for understanding and interpreting text rather than generating it.

### How BERT Differs from GPT-3 in Text Generation

- BERT is an encoder-only model, focusing on **understanding text**, while GPT-3 is a decoder-only model designed for **generating text**.
- BERT uses **bidirectional context**, which makes it better for tasks like classification, question answering, and sentence similarity.
- GPT-3 uses **left-to-right (autoregressive) context**, which makes it better for producing fluent and coherent text.

## 4. Prompt Engineering and Control of Output:

- **Q12:** Write a prompt for a language model to generate a 150-word description of a futuristic city. Explain the role of clarity and specificity in the prompt.

**Answer:**

**Example Prompt:**

"Write a 150-word descriptive paragraph about a futuristic city in the year 2150. The city should feature advanced transportation systems like flying cars and magnetic trains, sustainable energy sources such as solar panels and wind turbines, and smart infrastructure powered by AI. Include details about the appearance of buildings, the lifestyle of the residents, and how technology shapes daily life."

### Role of Clarity and Specificity in the Prompt:

Clarity ensures that the model understands what is expected. By clearly stating the required length (150 words) and purpose (description), the model avoids producing irrelevant or overly brief content.

Specificity guides the model on **what to include** (e.g., flying cars, AI-powered infrastructure, lifestyle of residents). It reduces ambiguity, helping the model stay focused and generate richer, more accurate content. A vague prompt like "Describe a city" might result in a generic or unrelated response, whereas a detailed prompt leads to a more useful and targeted output.

- **Q13:** How can **temperature** and **max tokens** be adjusted in a language generation model to control the creativity and length of the generated output? Provide examples of both adjustments



### Answer:

#### Temperature

Temperature controls the randomness or creativity of the text generated by a language model.

- **Low temperature (e.g., 0.2):** Makes the output more focused, predictable, and conservative. It chooses the most likely next words.
- **High temperature (e.g., 0.9):** Increases randomness, allowing more creative or unexpected word choices.

#### Example:

Prompt: "The robot walked into the room and..."

- *Temperature 0.2:*  
"The robot walked into the room and sat down quietly in the corner."

#### Max Tokens

Max tokens set the maximum number of words or characters the model is allowed to generate.

- **Low max tokens (e.g., 50):** Output will be short and may cut off abruptly.
- **High max tokens (e.g., 300):** Output will be longer, allowing for more detailed or complete responses.

#### Example:

Prompt: "Describe a futuristic vehicle."

- *Max tokens 50:*  
"The vehicle is sleek, fast, and powered by electricity. It uses sensors..."

- **Q14:** Write a prompt to generate a dialogue between two characters in a mystery novel. Provide guidelines in your prompt for tone and character development.

### Answers:

#### Example Prompt:

"Generate a dialogue between Detective Elara, a calm and methodical investigator in her 40s, and Jordan, a nervous young suspect in his 20s who claims to have witnessed a crime. The setting is a quiet interrogation room at night. The tone should be tense and suspenseful. Elara should ask probing questions while remaining composed. Jordan should sound anxious, stuttering occasionally, and revealing key details slowly. Use vivid language to show their emotions and hint at a deeper mystery without revealing the full truth yet. The dialogue should be 200 words."

#### Explanation of Guidelines:

- **Tone:** "Tense and suspenseful" helps the model keep the mystery alive.
- **Character Development:** Clear descriptions of Elara and Jordan guide the model to write distinct voices and emotions.
- **Setting & Mood:** Mentioning the interrogation room and time of day adds atmosphere.
- **Word Limit:** Keeps the output concise and focused on meaningful dialogue.

## 5. Evaluating the Output of Generative AI Models:

- **Q15:** How would you evaluate the quality of text generated by a model like GPT-3? List at least three criteria you would consider when assessing its output.

**Answer:**

### 1. Relevance to the Prompt

The output should directly respond to the given prompt or question. It must stay on topic and address the user's intent without going off-track or introducing unrelated content.

*Example:* If the prompt asks for a summary of climate change causes, the model shouldn't talk about space exploration.

### 2. Coherence and Fluency

The text should be logically organized, grammatically correct, and easy to read. Ideas should flow smoothly from one sentence to the next, forming a coherent narrative or explanation.

*Example:* Sentences should not contradict each other or seem randomly placed.

### 3. Factual Accuracy

The information presented should be correct and verifiable, especially in tasks that involve real-world knowledge. Outdated or misleading statements reduce the value of the output.

*Example:* Saying "Mars has breathable air" would be a clear factual error.

- **Q16:** What are some common problems with generated content, such as **hallucinations** or **irrelevant responses**? How can these issues be minimized in prompt design?

**Answer:**

### Common Problems with Generated Content

### 1. **Hallucinations**

The model generates content that is factually incorrect or made-up. For example, it might invent events, statistics, or quotes that don't exist.

### 2. **Irrelevant Responses**

The response may go off-topic or provide information unrelated to the prompt. This usually happens when the prompt is vague or ambiguous.

### 3. **Repetition**

The model may repeat the same idea or phrase unnecessarily, especially in longer outputs.

### 4. **Overconfidence in Incorrect Answers**

The model may state false information confidently, which can mislead users if unchecked.

## **How to Minimize These Issues with Prompt Design**

### 1. **Be Clear and Specific**

Clearly state what you want in the prompt. For example, instead of "Tell me about planets," ask "List the names of all eight planets in our solar system with one fact each."

### 2. **Set Constraints**

Define format, tone, length, or structure. For example, "Write a 100-word paragraph explaining how solar panels work, in simple language."

### 3. **Avoid Ambiguity**

Use direct language and avoid terms that could be interpreted in multiple ways.

### 4. **Include Context**

Provide background information if needed so the model has enough reference to stay on topic.

- **Q17: How can **feedback loops** be used to improve generative models? Explain how iterative testing and refinement of prompts can enhance the output.**

### **Answer:**

A feedback loop involves evaluating the model's output, analyzing its quality, and then using that analysis to make improvements—either in how the model is prompted or how it is fine-tuned. This loop continues until the generated content meets the desired standard.

## **How Feedback Improves Generative Models**

### 1. **Prompt Refinement**

If a prompt produces vague or off-topic results, you can modify the prompt to be

more specific or structured. Over multiple attempts, this improves the clarity of responses.

*Example:*

Original prompt: "Explain climate change."

Improved prompt: "Explain the main causes of climate change in under 100 words using simple language."

## 2. **Error Correction**

Feedback can help identify hallucinations, irrelevant content, or poor grammar in the output. This allows you to fine-tune the model or adjust input to reduce such issues.

## 3. **Task Optimization**

By repeatedly testing prompts and observing the responses, you can discover the best structure and wording for certain types of tasks like summarization, translation, or question answering.

## 4. **Training Signals (in fine-tuning)**

In more advanced workflows, feedback can be used as training data. High-quality outputs are reinforced, and low-quality ones are penalized, guiding future generations.

# 6. 7. Hands-on Practice with Generative AI:

- **Q18:** Write a prompt that will instruct a language model to summarize a research paper about machine learning. Include specific instructions to highlight the main points and avoid irrelevant details.

**Answer:**

**Example Prompt:**

"Summarize the following research paper on machine learning in 150–200 words. Focus on the paper's objective, key methods or algorithms used, main findings, and potential applications. Avoid mentioning author names, citations, or background details that are not central to the research. Keep the summary concise, objective, and informative."

**Explanation:**

This prompt includes:

- **Clear task:** Summarize a research paper
- **Length control:** 150–200 words

- **Focus areas:** Objective, methods, results, applications
- **Exclusions:** Author names, citations, unnecessary background
- **Tone/style:** Concise and informative

Such specificity helps the model generate a focused, accurate summary that's useful for readers trying to quickly understand the core ideas of the research.

- **Q19:** Generate a list of ideas for a new mobile app using a language generation model. Provide at least five app ideas and explain how the model can generate creative suggestions.

**Answer:**

#### App Ideas Generated by a Language Model:

##### 1. **EcoHabit**

An app that helps users build eco-friendly habits by tracking their daily actions (like saving water, reducing plastic use) and rewarding them with badges and challenges.

##### 2. **MindMap Journal**

A visual journaling app that uses AI to convert your written thoughts into colorful mind maps, helping with self-reflection, goal tracking, and mental clarity.

##### 3. **MoodChef**

A recipe suggestion app that recommends meals based on the user's mood, weather, and available ingredients at home.

##### 4. **SkillSwap**

A peer-to-peer learning app where users exchange skills (e.g., photography for language lessons) by scheduling video sessions or local meetups.

##### 5. **StorySprout**

A creative writing app for kids and teens that uses generative AI to help them build characters, plots, and endings based on their chosen themes or genres.

#### How the Model Generates Creative Suggestions:

- **Trained on Diverse Data:** Language models are trained on books, websites, app stores, and user interactions, which give them a broad view of existing ideas and gaps.
- **Pattern Recognition:** The model identifies what makes an idea unique, feasible, and interesting.
- **Combination of Concepts:** It combines unrelated concepts in novel ways (e.g., mood + food, skills + community).

- **Prompt-Driven Creativity:** With the right prompts (like "generate 5 creative mobile app ideas"), the model outputs diverse and imaginative ideas tailored to the request.
- **Q21:** Generate a set of **product descriptions** for an e-commerce website using a language model. Evaluate the clarity, persuasiveness, and accuracy of the descriptions.

**Answer:**

### Example Product Descriptions

#### 1. **Wireless Noise-Cancelling Headphones**

Experience immersive sound with our premium wireless headphones. Featuring active noise cancellation, 30-hour battery life, and ultra-soft ear cushions, these headphones are perfect for travel, work, or relaxing at home. Bluetooth 5.2 ensures a stable and fast connection for uninterrupted listening.

#### 2. **Smart Water Bottle with Hydration Tracker**

Stay healthy and hydrated with this smart water bottle. It tracks your water intake throughout the day and sends reminders to your phone. Made with BPA-free materials, it's leak-proof, rechargeable, and available in 5 vibrant colors.

#### 3. **Ergonomic Office Chair**

Upgrade your workspace with an ergonomic chair designed for comfort and productivity. Adjustable lumbar support, breathable mesh back, and 360° swivel make it ideal for long workdays. Holds up to 125 kg and suits all desk heights.

### Evaluation of the Descriptions

- **Clarity:**  
Each description is written in simple, direct language. Features and benefits are easy to understand.
- **Persuasiveness:**  
The descriptions highlight unique selling points (e.g., "30-hour battery," "hydration reminders," "adjustable lumbar support") to appeal to customer needs and emotions.
- **Accuracy:**  
Assuming the features match the actual products, the descriptions are specific and accurate. However, factual accuracy depends on the real product specs — always verify with actual data before use.